

Analysis of the Top Three Frontier AI Models

Opus 5.5 leads on coding and analyst work, Astra on 3D tools and hard science, and Kimi K3 on cost and self-hosting



Claude Opus 5.5

Anthropic | Released Sep 22, 2026

- Best for **coding, agent workflows and analyst work** like research and modeling
- Scored **66.4% on Terminal-Bench 4.0** vs. 57.9% for Astra; one customer migrated **680K lines of code in a day**
- Costs **\$4 / \$20 per 1M tokens**, about 60% less than Astra



GPT-6 Astra

OpenAI | Released Sep 3, 2026

- Best for **operating desktop software, 3D tools like Blender and UE5**, and hard math and science
- Scored **97.6% on FrontierMath Tier 4** and 72.6% on OSWorld 2.0
- Most expensive at **\$10 / \$50 per 1M tokens**, with access still rolling out in stages



Kimi K3

Moonshot AI | Released Jul 16, 2026

- Best for **high-volume, low-cost work** and anything that must be **self-hosted**
- The **only open-weight model** of the three, with 2.8T parameters
- Cheapest at **\$3 / \$15 per 1M tokens**, but scores 44 vs. 58 for Opus on the Artificial Analysis index

Key Takeaways

- **Opus 5.5** should be the default for most work, given its lead on coding and knowledge-work benchmarks and its **lower price than Astra**
- **Astra** is worth the premium for specialist work, while **Kimi K3** fits teams that need to keep data in-house or cut costs